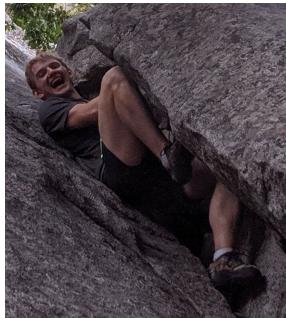
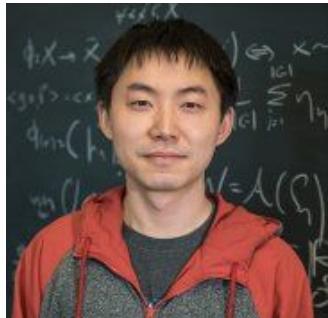




Memorization in Language Models

Katherine Lee
Cornell, Google Brain, May 13, 2022





Katherine Lee

Daphne Ippolito

Nicholas Carlini

Chiyuan Zhang

Matthew
Jagielski

Florian Tramèr



Andrew Nystrom

David Mimno

Hannah Brown

Fatemehsadat
Mireshghallah

Reza Shokri

Language Model

Given some preceding text, predict what the next word is

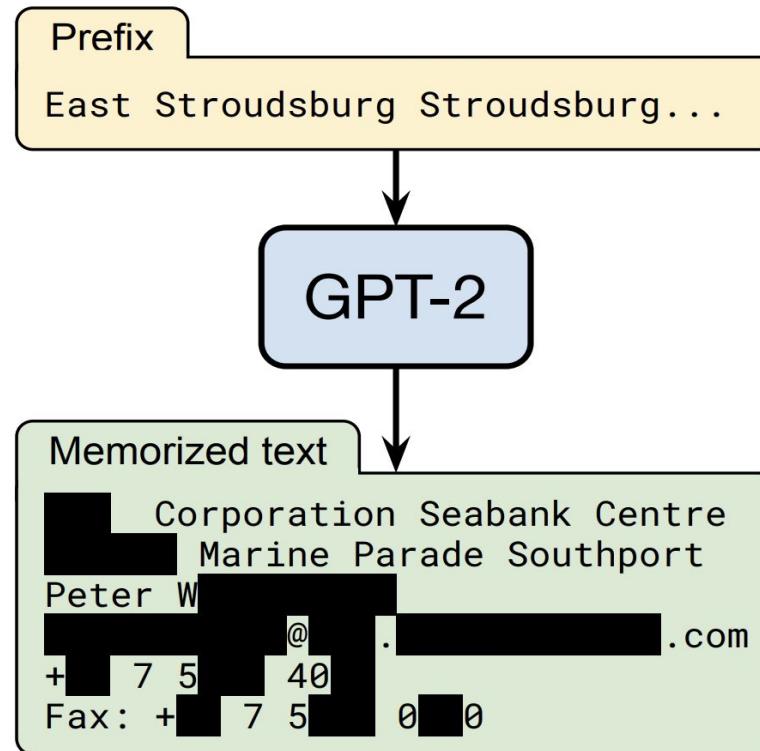
$$P(\text{word} \mid \text{all previous words})$$

Large Models are Leaky

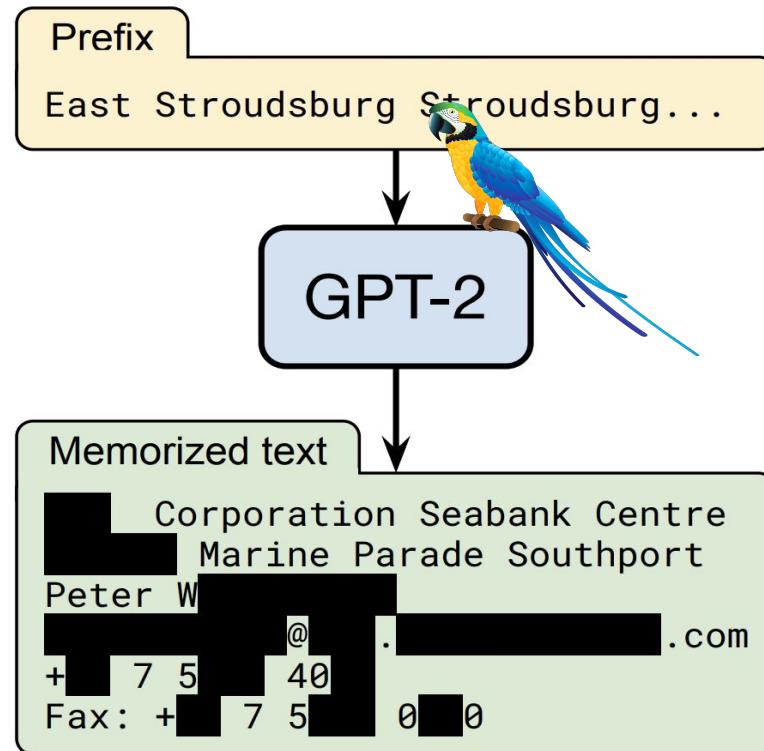


WHEN YOU TRAIN PREDICTIVE MODELS
ON INPUT FROM YOUR USERS, IT CAN
LEAK INFORMATION IN UNEXPECTED WAYS.

Large Models are Leaky

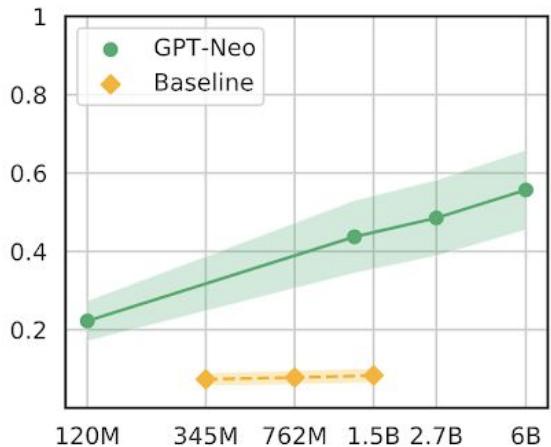


Large Models are Leaky



Discoverable memorization scales...

...with model scale



(a) Model scale

Neural language models memorize training data.

Our fully equipped family sized lodges offer a comfortable luxurious stay for a fantastic price, giving you beautiful views of the lakes and the surrounding countryside. Offering luxurious self-catering holidays in our fully featured Scandinavian holiday lodges. Perfectly located to explore the beaches, coastline.

you are only looking to find rent to own homes in your city or are open to exploring all kinds of rent to own home listings, our database does it all. One of the best aspects of iRentToOwn.com is that, besides options to rent to buy a house, it has numerous other categories of home sale options. These include bank foreclosure homes, pre-foreclosure homes, short sales, HUD/government foreclosures, auction homes and owner-financing/FSBO (For Sale By Owner) homes. With help from the convenient search features offered by our site, shoppers are able to find their ideal lease to own home, real estate company, and more

you'll need to be knowledgeable to make the very best decisions. We will make sure you know what can be expected. We take the surprises from the picture by giving accurate and thorough information. You can start by talking about your task with our client service staff when you dial 888-353-1299. We'll address all of your questions and arrange the initial meeting. We work closely with you through the whole project, and our team can show up promptly and prepared

Our journeyman also does service work, troubleshooting when a breaker fails or a light stops working. Our journeyman does not offer permits that must be issued by our master. Our journeyman follows our master's plans and directions. Our journeyman's responsibilities will vary based on the work that needs to be done. Our journeymen are skilled with residential, commercial and industrial installations and repairs.ust work from six years as an apprentice, under direct supervision of our master, and pass a journeyman test.

Neural language models memorize training data.

Our fully equipped family sized lodges offer a comfortable luxurious stay for a fantastic price, giving you beautiful views of the lakes and surrounding countryside.

Offering luxurious self-catering lodges in our fully featured Scandinavian holiday lodges. Perfectly located to explore the beaches, coastline.

571x

you are only looking to find rent to own homes in your city or are open to exploring all kinds of rent to own home listings, our database does it all. One of the best aspects of iRentToOwn.com is that, besides options to rent to buy a house, it has numerous other categories of home sale options. These include bank foreclosure homes,

51x
pre-foreclosure homes, short sales, HUD/government foreclosures, auction homes and owner-financing/FSBO (For Sale By Owner) homes. With help from the convenient search features offered by our site, shoppers are able to find their ideal lease to own home, real estate company, and more

you'll need to be knowledgeable to make the very best decisions. We will make sure you know what can be expected. We take the surprises from the picture by giving accurate and thorough information. You can start by talking about your task with our customer service staff when you dial 888-353-1299. We'll address all of your questions and arrange the initial meeting. We work closely with you through the whole project, and our team can show up promptly and prepared

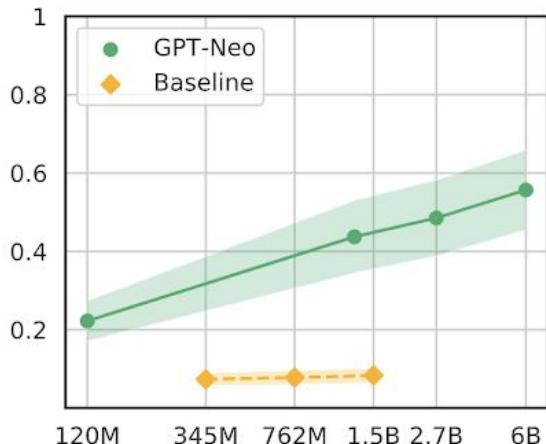
5,497x

Our journeyman also does service work, troubleshooting when a breaker fails or a light stops working. Our journeyman does not offer permits that must be issued by our master. Our journeyman follows our master's plans and directions. Our journeyman's responsibilities will vary based on the work that needs to be done. Our journeymen are skilled with residential, commercial and industrial installations and repairs. Just work from six years as an apprentice, under direct supervision of our master, and pass a journeyman test.

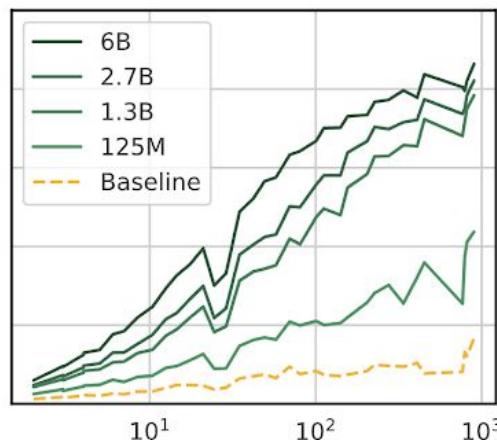
6x

Discoverable memorization scales...

...with data repetition

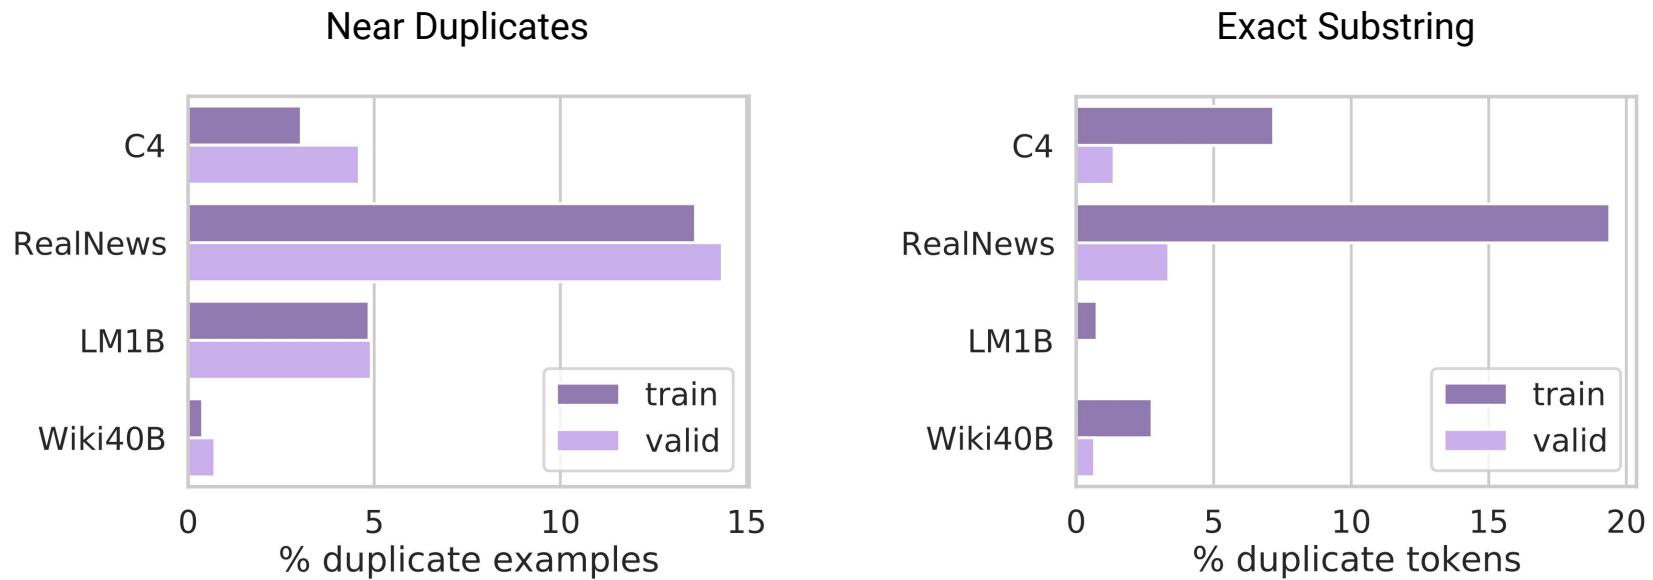


(a) Model scale



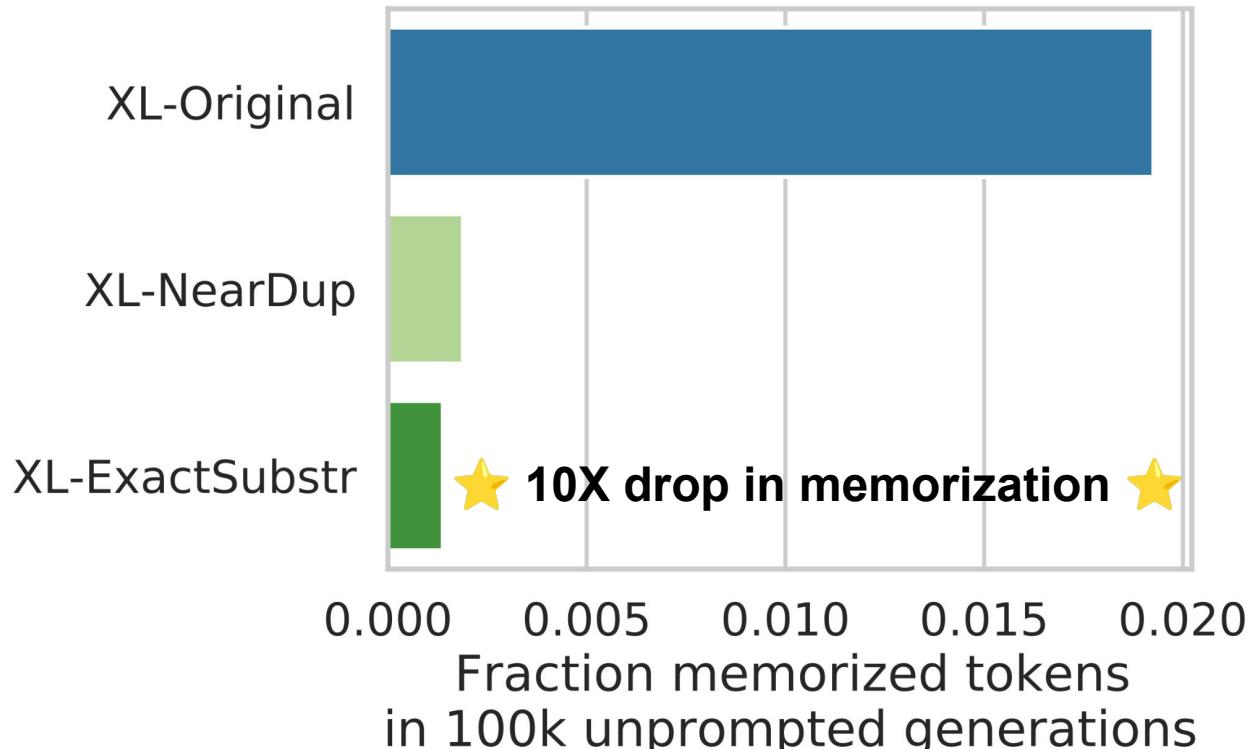
(b) Data repetition

Deduplicating Text Data

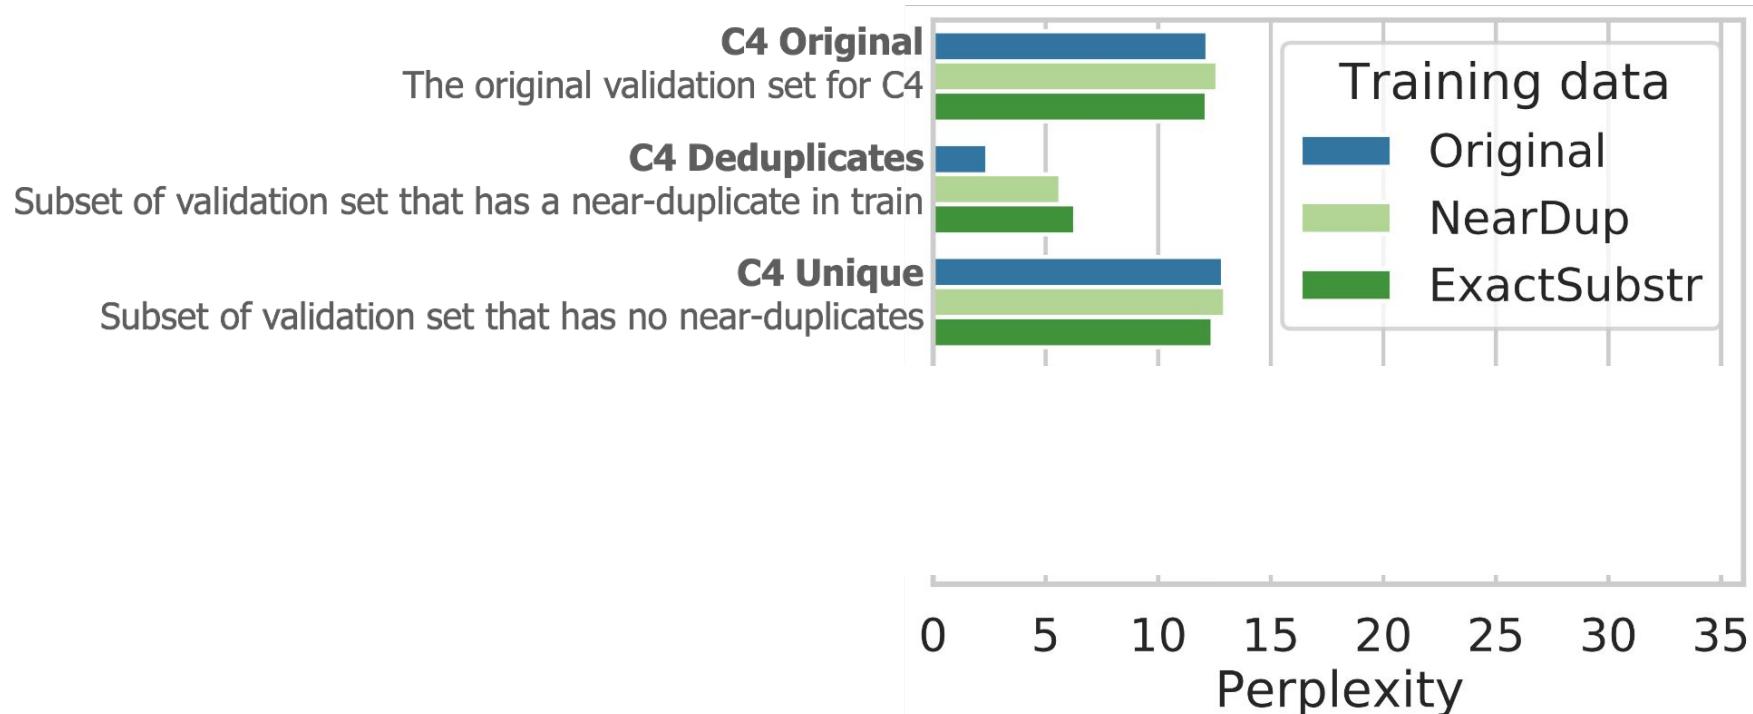


Generate text
with random
sampling
(top-k=50)

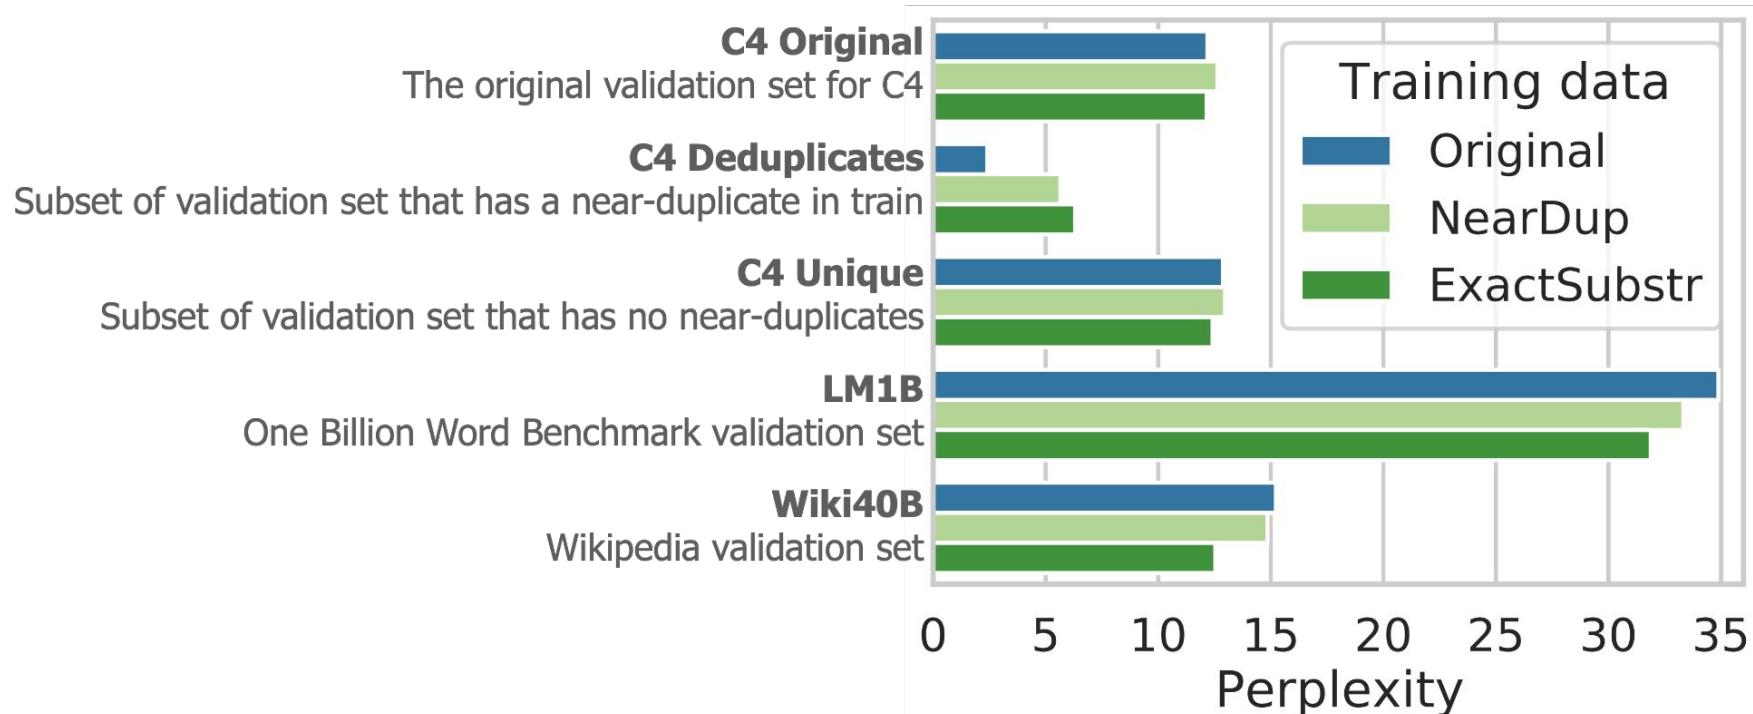
Unprompted Memorization



Deduplicated models are better.



Deduplicated models are better.



Neural language models memorize training data.

Our fully equipped family sized lodges offer a comfortable luxurious stay for a fantastic price, giving you beautiful views of the lakes and the surrounding countryside. Offering luxurious self-catering holidays in our fully featured Scandinavian holiday lodges. Perfectly located to explore the beaches, coastline.

you are only looking to find rent to own homes in your city or are open to exploring all kinds of rent to own home listings, our database does it all. One of the best aspects of iRentToOwn.com is that, besides options to rent to buy a house, it has numerous other categories of home sale options. These include bank foreclosure homes, pre-foreclosure homes, short sales, HUD/government foreclosures, auction homes and owner-financing/FSBO (For Sale By Owner) homes. With help from the convenient search features offered by our site, shoppers are able to find their ideal lease to own home, real estate company, and more

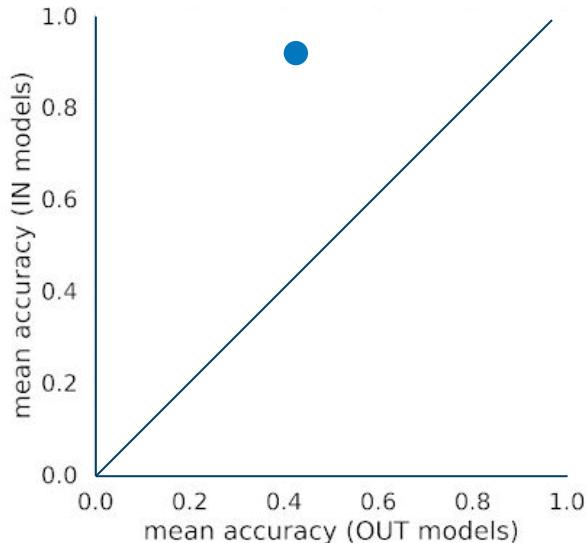
you'll need to be knowledgeable to make the very best decisions. We will make sure you know what can be expected. We take the surprises from the picture by giving accurate and thorough information. You can start by talking about your task with our client service staff when you dial 888-353-1299. We'll address all of your questions and arrange the initial meeting. We work closely with you through the whole project, and our team can show up promptly and prepared

Our journeyman also does service work, troubleshooting when a breaker fails or a light stops working. Our journeyman does not offer permits that must be issued by our master. Our journeyman follows our master's plans and directions. Our journeyman's responsibilities will vary based on the work that needs to be done. Our journeymen are skilled with residential, commercial and industrial installations and repairs.ust work from six years as an apprentice, under direct supervision of our master, and pass a journeyman test.

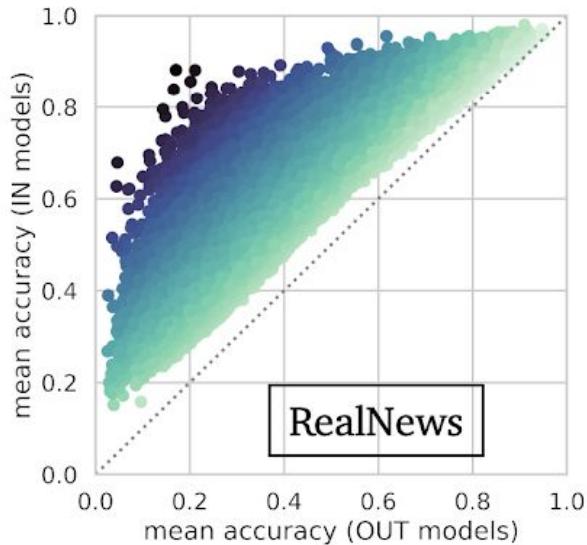
Counterfactual Memorization

Which examples are rarely seen but still memorized?

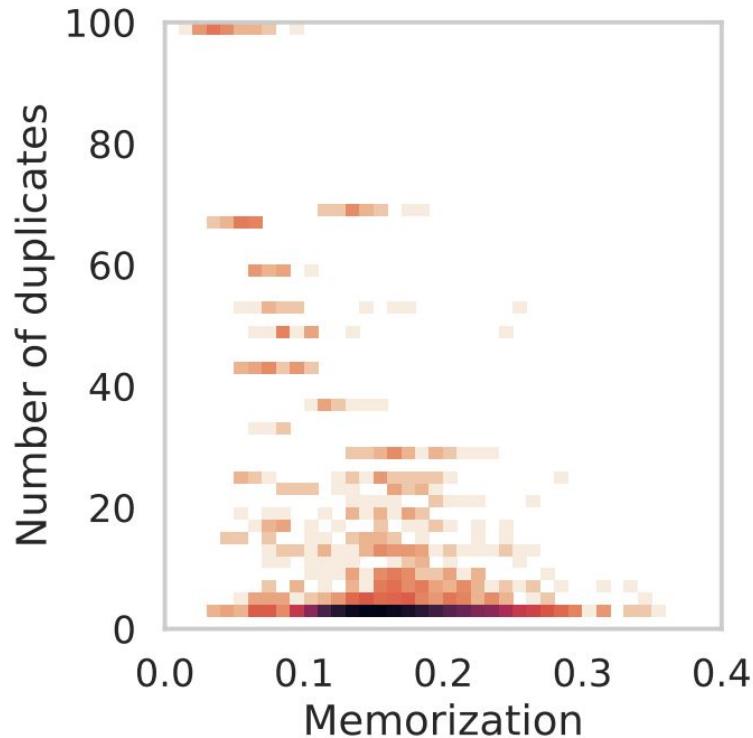
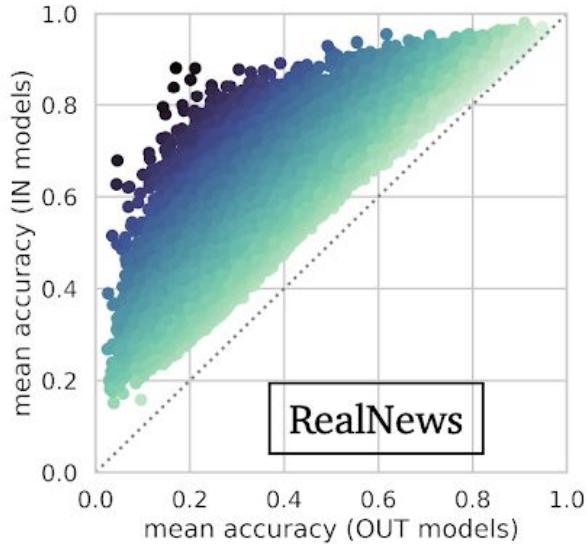
Counterfactual Memorization



Counterfactual Memorization



Counterfactual Memorization

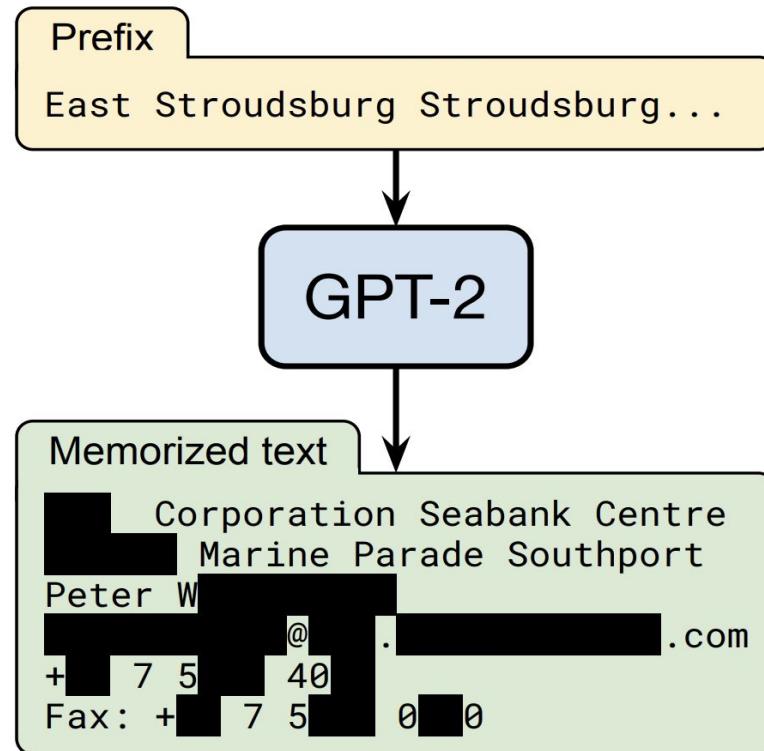


Prompting

Given some preceding text, predict what the next word is

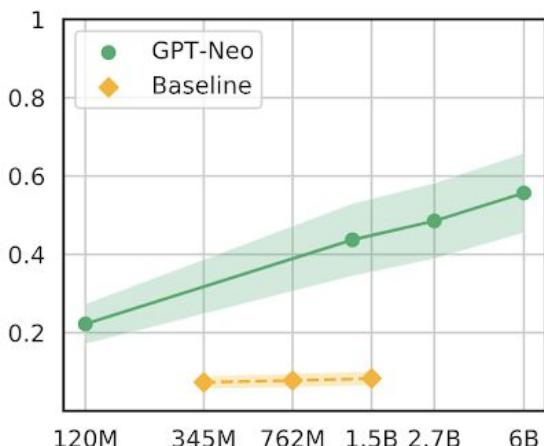
$$P(\text{word} \mid \text{all previous words, prompt})$$

Large Models are Leaky

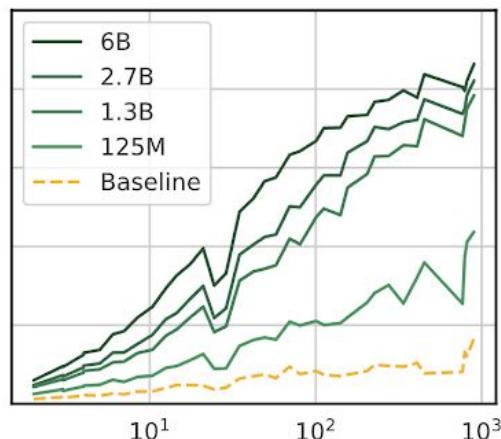


Discoverable memorization scales...

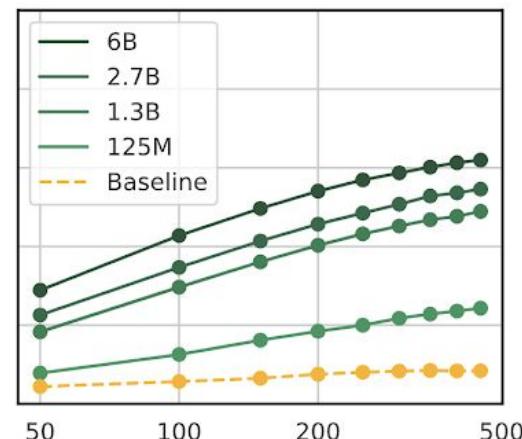
...with context length



(a) Model scale



(b) Data repetition



(c) Context size

Types of Memorization

Private / sensitive information: "My social security number is XXXXX."

Common phrases: "To whom it may concern..."

Facts: "Christmas is celebrated on Dec 25th"

Quotes: "Trump said, 'Tariffs are the greatest!'"

Quotes: "Sally Smith said, 'Sam is the worst.'"

What Does it Mean for a Language Model to Preserve Privacy?

How do we use language?



How do we use language?

Communicate ideas, wants, needs

Form identity and community

Expression



Privacy concerns
are as broad as
those of real life

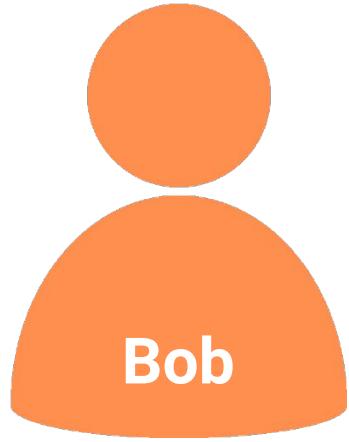
Language is contextual

Shared information ≠ public information

Information may be private to only some people

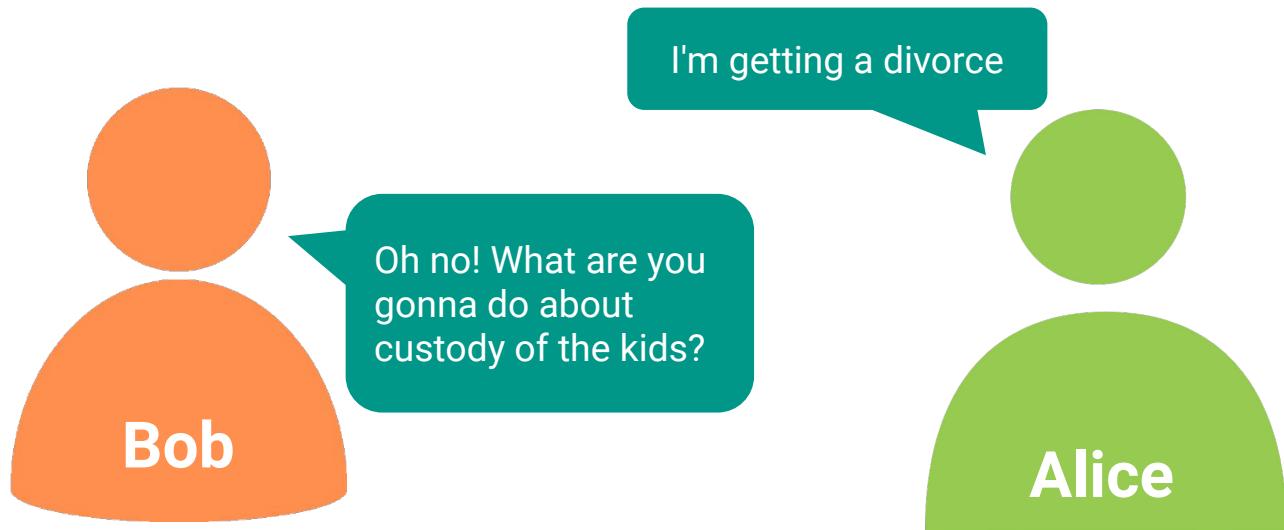
Or only in some contexts

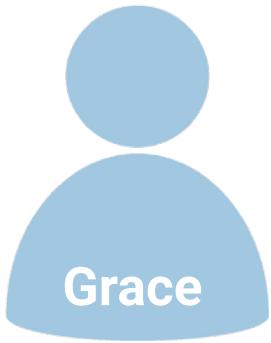
Identifying all of this is hard!



I'm getting a divorce





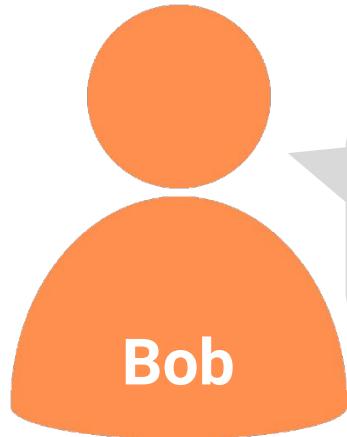


Grace



Charlie

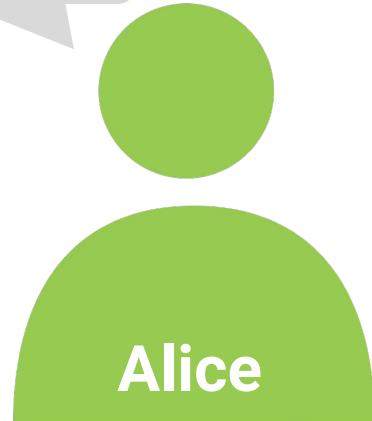
Did you hear Alice
is getting divorced?



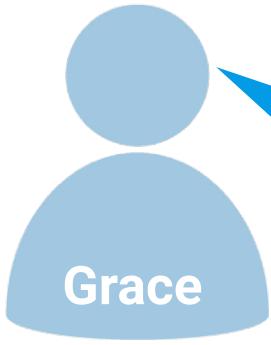
Bob

I'm getting a divorce

Oh no! What are you
gonna do about
custody of the kids?



Alice



Grace

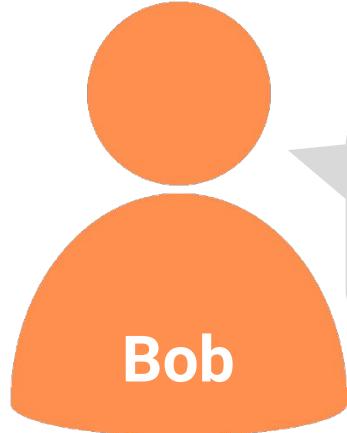
Wait Alice is getting a divorce?



Charlie

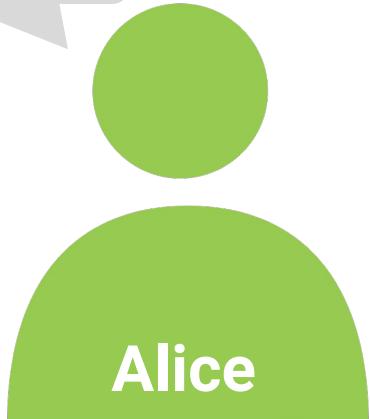
No way!

Did you hear Alice is getting divorced?



Bob

Oh no! What are you gonna do about custody of the kids?



Alice

I'm getting a divorce

We *memorize*
information

We *memorize*
information
then judge the
context

We memorize
information
then judge the
context

BUT

Language models don't
have this understanding!

Information for context
usually is beyond data
given

Privacy is not binary

Privacy violations range in severity

When is sharing okay?

Who can we share with?

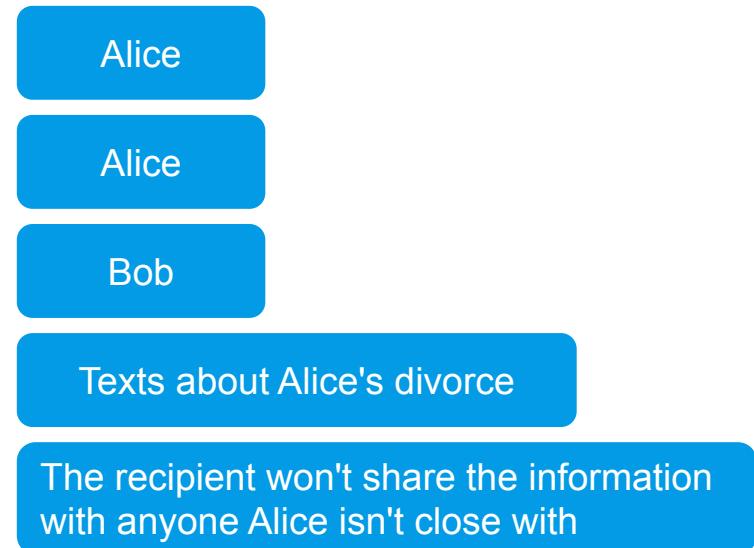
What is the private information?

All heavily context dependent and can change



Contextual Integrity

- 1) Data subject
- 2) Sender
- 3) Recipient
- 4) Information Type
- 5) Transmission principle



Current NLP Privacy Methods

Why can't we just remove private text?

[aka, text sanitization]

Private information has no one format

lastname AT website DOT com

Language constantly changes

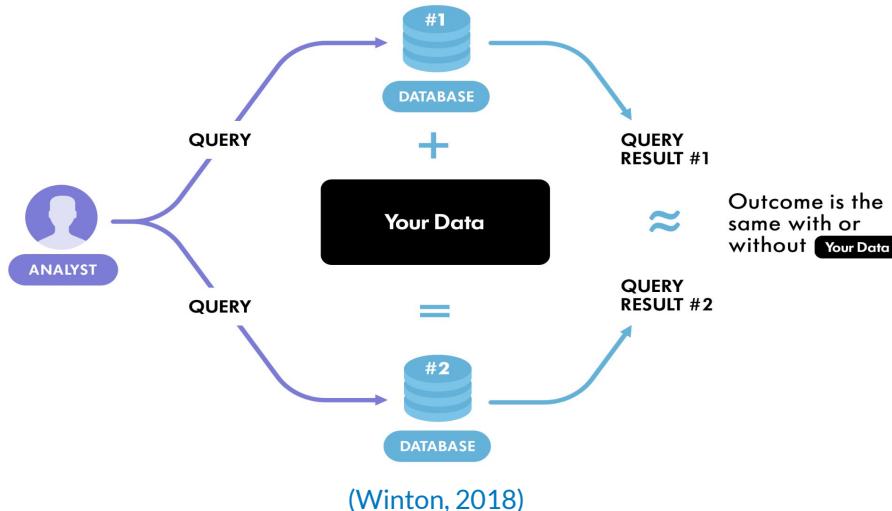
Die → Unalive

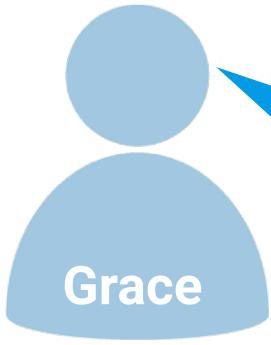
Privacy is context dependent

We're throwing Bob a surprise party!

What about differential privacy?

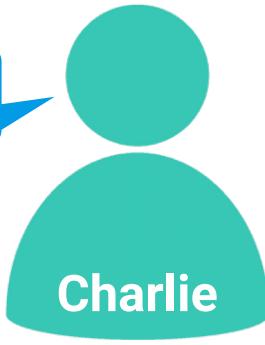
For some value ϵ , and algorithm A, the probability of a single record being in the training dataset of A is indistinguishable (*relative to ϵ*) from the probability that it is not (Dwork, 2006).





Grace

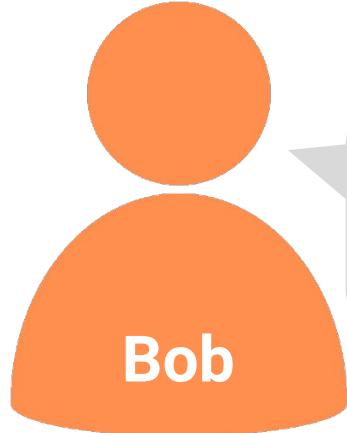
Wait Alice is getting a divorce?



Charlie

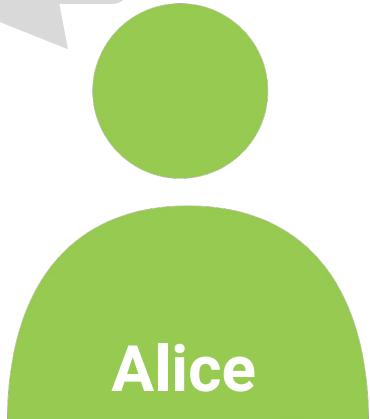
No way!

Did you hear Alice is getting divorced?



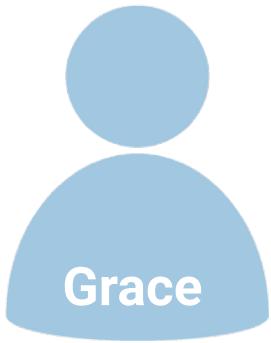
Bob

Oh no! What are you gonna do about custody of the kids?



Alice

I'm getting a divorce

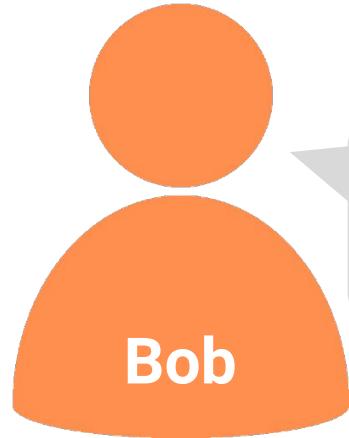


Grace



Charlie

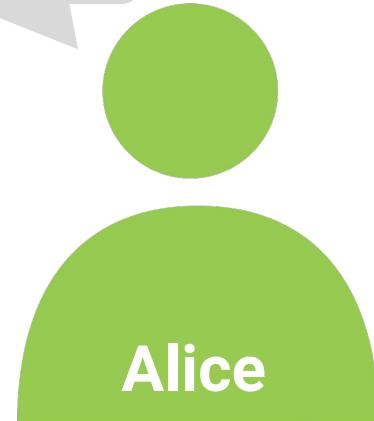
Did you hear Alice is
in a custody battle?



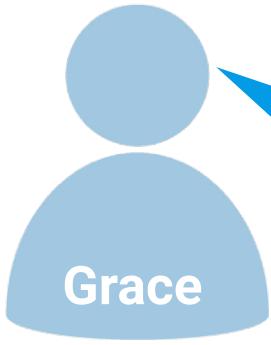
Bob

I'm getting a divorce

Oh no! What are you
gonna do about
custody of the kids?



Alice



Does that mean
she's getting
divorced?



No way!

Did you hear Alice is
in a custody battle?

A large orange circular icon representing a person's head and shoulders. The name "Bob" is written in white capital letters at the bottom.

I'm getting a divorce

A large green circular icon representing a person's head and shoulders. The name "Alice" is written in white capital letters at the bottom.

Oh no! What are you
gonna do about
custody of the kids?

A grey rounded rectangular speech bubble containing text.

Can shared information be private?



The Panama Papers: Exposing the Rogue Offshore Finance Industry

(ICIJ, 2016)

DP makes assumptions

Privacy is ***binary***

Private information is ***identifiable***

Private information will ***never be shared***

Units of private information follow defined natural language units



DP makes assumptions

Privacy is ***binary***

Private information is ***identifiable***

Private information will ***never be shared***

Units of private information follow defined natural language units

Guarantees don't align with our ideas of privacy for language

Withholding any unit of data cannot guarantee privacy

Shared information gets less privacy guarantees

What is a record?

How can Language Models Preserve Privacy?

Can users consent?

One person's *data* includes multiple people's *information*

Privacy guarantees that do exist can't be easily explained

Informed consent is generally impossible



Publicly available ≠ publicly directed

Data can be shared without consent

Public posts on social media often have target audiences

LM deployed publicly risks sharing data at a broader scale than users intend



Privacy Preserving LMs?

Train on data intended to be public

Finetune locally on user-contributed
data if needed

Privacy is *meaningfully* preserved this
way



Questions & Thank you!

What violations of privacy do you accept?

Can informed consent be given?

What questions does this raise for researchers designing the technology?

What sort of data *should* we be using?

References

Carlini et al. Extracting Training Data from Large Language Models. USENIX SEC 2021.

Carlini, et al., Quantifying Memorization Across Neural Language Models, arxiv, 2022

Lee, et al. Deduplicating Training Data Makes Language Models Better, ACL 2022,

Zhang, et al., Counterfactual Memorization in Neural Language Models, arxiv, 2022

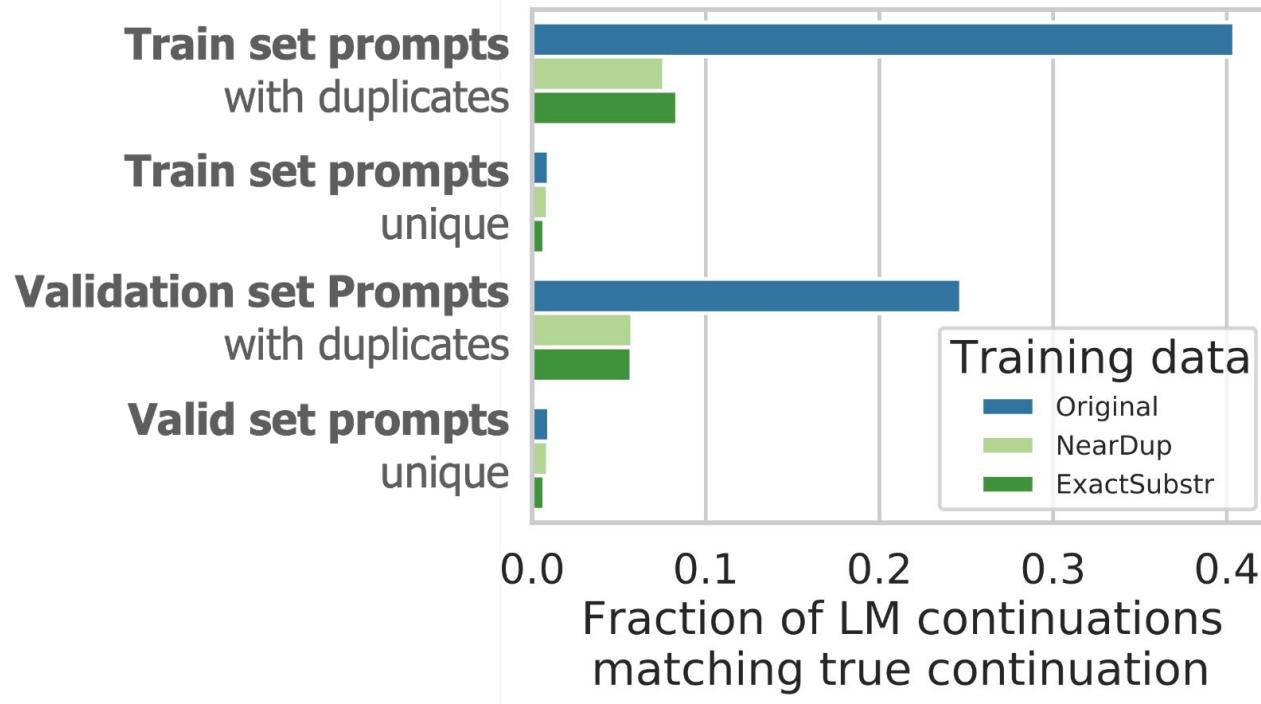
Brown, et al. “What Does it Mean for a Language Model to Preserve Privacy?” FAccT 2022

Thank you!

Extra slides

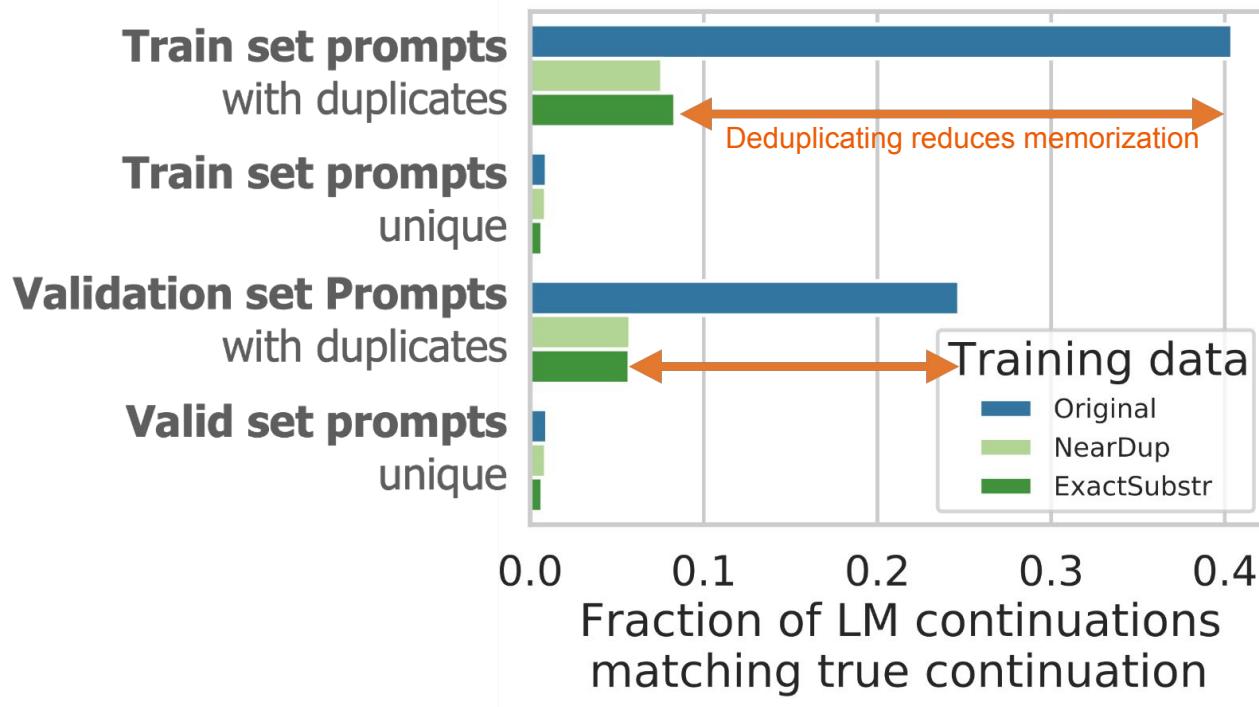
Prompt with
32 tokens,
generate with
top-k=50

Prompted Memorization



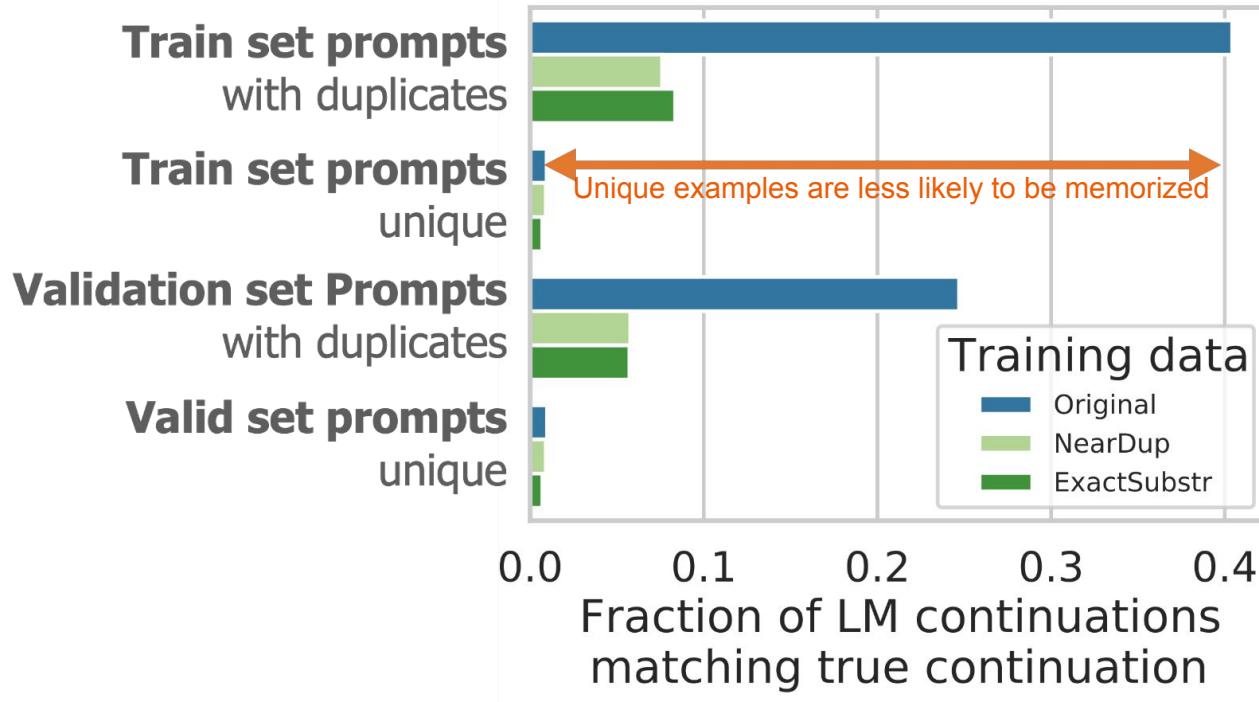
Prompt with
32 tokens,
generate with
top-k=50

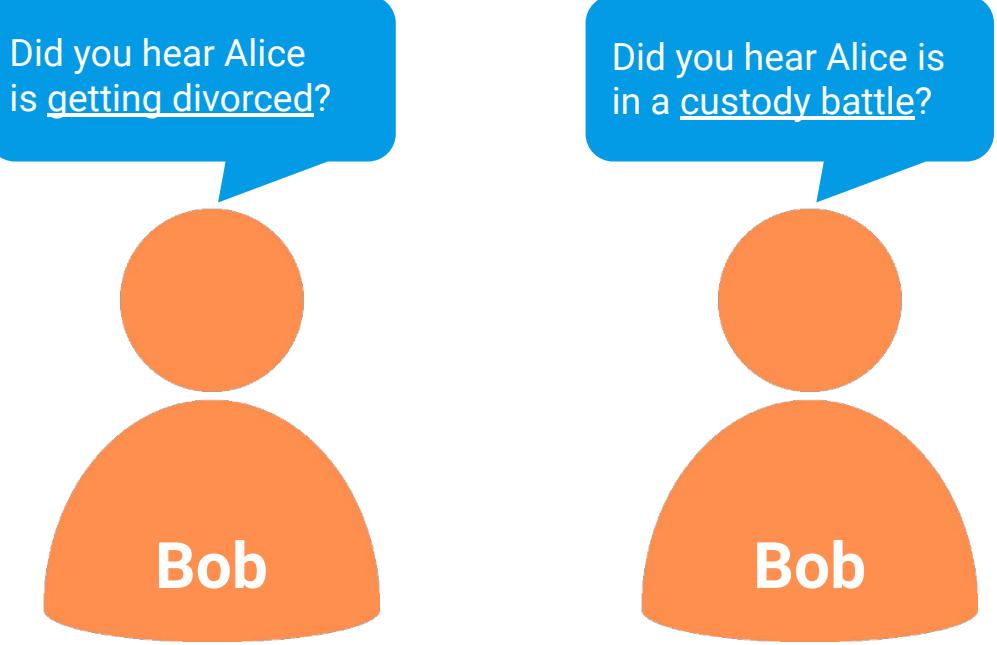
Prompted Memorization



Prompt with
32 tokens,
generate with
top-k=50

Prompted Memorization



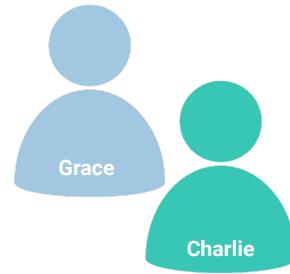


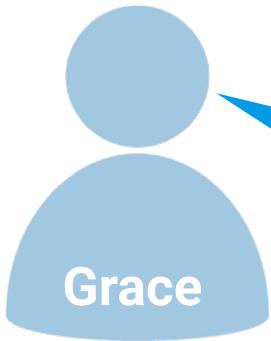
Did you hear Alice
is getting divorced?

Bob

Did you hear Alice is
in a custody battle?

Bob





Wait Alice is
getting divorced?

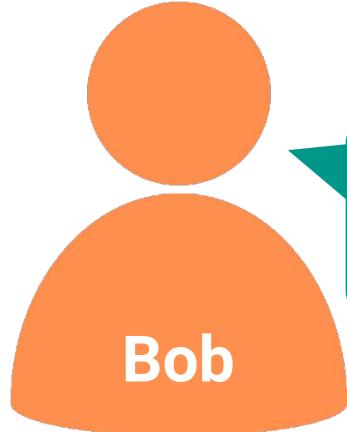


No way!

Did you hear Alice
is getting divorced?

A large orange circular icon representing a person, with the name "Bob" written in white capital letters below it. A large green circular icon representing a person, with the name "Alice" written in white capital letters below it. The two icons are positioned at opposite ends of a horizontal axis, facing each other.

I'm getting a divorce



Oh no! What are you
gonna do about
custody of the kids?



Who can private information be shared with?



Suicide hotline shares data with for-profit spinoff, raising ethical questions

(Levine, 2022)

References

- Hannah Brown, Katherine Lee, Fatemehsadat Mireshghallah, Reza Shokri, and Florian Tramèr, What Does it Mean for a Language Model to Preserve Privacy? *arXiv preprint arXiv:2202.05520*, 2022.
- Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. *arXiv preprint arXiv:2012.07805*, 2020.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.
- Cynthia Dwork, Frank McSherry, Kobbi Nissim, and Adam Smith. Calibrating noise to sensitivity in private data analysis. In *Theory of cryptography conference*, pages 265–284. Springer, 2006.
- International Consortium of Investigative Journalists. About the Panama Papers investigations.
<https://www.icij.org/investigations/panama-papers/pages/panama-papers-about-the-investigation/>, 2016.
- Alexandra S. Levine. Suicide hotline shares data with for-profit spinoff, raising ethical questions, Jan 2022.
- Helen Nissenbaum. *Privacy in context*. Stanford University Press, 2009.